

Machine Learning Worksheet 6

Linear Classification

1 Linear separability

Problem 1: Given a set of data points $\{\mathbf{x}_n\}$, we can define the *convex hull* to be the set of all points \mathbf{x} given by

$$\mathbf{x} = \sum_n \alpha_n \mathbf{x}_n$$

where $\alpha_n \geq 0$ and $\sum_n \alpha_n = 1$. Consider a second set of points $\{\mathbf{y}_n\}$ together with their corresponding convex hull. By definition, the two sets of points will be linearly separable if there exists a vector \mathbf{w} and a scalar w_0 such that $\mathbf{w}^T \mathbf{x}_n + w_0 > 0$ for all \mathbf{x}_n , and $\mathbf{w}^T \mathbf{y}_n + w_0 < 0$ for all \mathbf{y}_n . Show that if their convex hulls intersect, the two sets of points cannot be linearly separable, and conversely that if they are linearly separable, their convex hulls do not intersect.

Problem 2: Show that for a linearly separable data set, the maximum likelihood solution for the logistic regression model is obtained by finding a vector \mathbf{w} whose decision boundary $\mathbf{w}^T \phi(\mathbf{x}) = 0$ separates the classes and then taking the magnitude of \mathbf{w} to infinity.

2 Multiclass classification

Problem 3: Consider a generative classification model for K classes defined by prior class probabilities $p(\mathcal{C}_k) = \pi_k$ and general class-conditional densities $p(\phi|\mathcal{C}_k)$ where ϕ is the input feature vector. Suppose we are given a training data set $\{\phi_n, t_n\}$ where $n = 1, \dots, N$, and t_n is a binary target vector of length K that uses the 1-of- K coding scheme, so that it has components $t_{nj} = I_{jk}$ if pattern n is from class \mathcal{C}_k . Assuming that the data points are drawn independently from this model, show that the maximum-likelihood solution for the prior probabilities is given by $\pi_k = \frac{N_k}{N}$ where N_k is the number of data points assigned to class \mathcal{C}_k .

3 Bounds

Problem 4: Suppose we test a classification method on a set of n new test cases. Let $X_i = 1$ if the classification is wrong and $X_i = 0$ if it is correct. Then $\hat{X} = n^{-1} \sum X_i$ is the observed error rate. If we regard each X_i as a Bernoulli with unknown mean p , then p should be the true, but unknown, error rate of our method. How likely is \hat{X} to not be within ε of p . How many test cases are necessary to ensure that the observed error rate is with probability at most 5% farther than 0.01 away from the true one?

4 The perceptron ★

An important example of a so called *linear discriminant* model is the *perceptron* of Rosenblatt. The following questions will look more closely at this algorithm. We will assume the following:

- The parameters of the perceptron learning algorithm are called *weights* and are denoted by \mathbf{w} .
- The *training set* consists of training inputs \mathbf{x}_i with labels $t_i \in \{+1, -1\}$.
- The *learning rate* is 1.
- Let k denote the number of weight updates the algorithm has performed at some point in time and \mathbf{w}^k the weight vector after k updates (initially, $k = 0$ and $\mathbf{w}^0 = \mathbf{0}$).
- All training inputs have bounded euclidean norms, i.e. $\|\mathbf{x}_i\| < R$, for all i and some $R \in \mathbb{R}^+$.
- There is some $\gamma > 0$ such that $t_i \tilde{\mathbf{w}}^T \mathbf{x}_i > \gamma$ for all i and some suitable $\tilde{\mathbf{w}}$ (γ is called a *finite margin*).

Problem 5: Write down the perceptron learning algorithm.

Problem 6: Given the following training set \mathcal{D} of labeled 2D training inputs, find a *separating hyperplane* using the perceptron learning rule. Illustrate the consecutive updates of the weight \mathbf{w} with a series of plots (do not plot the bias weight)!

$$\begin{aligned} \mathcal{D} = & \{((-0.7, 0.8), +1), ((-0.9, 0.6), +1), ((-0.3, -0.2), +1), ((-0.6, 0.7), +1)\} \\ & \cup \{((0.6, -0.8), -1), ((0.2, -0.5), -1), ((0.3, 0.2), -1)\} \end{aligned}$$

You will now show that the perceptron algorithm converges in a finite number of updates (if the training data is linearly separable).

Problem 7: Let \mathbf{w}^k be the k^{th} update of the weight during the perceptron algorithm. Show that $(\tilde{\mathbf{w}}^T \mathbf{w}^k) \geq k\gamma$. (Hint: How are $(\tilde{\mathbf{w}}^T \mathbf{w}^k)$ and $(\tilde{\mathbf{w}}^T \mathbf{w}^{k-1})$ related?)

Problem 8: Show that $\|\mathbf{w}^k\|^2 < kR^2$. Note that the algorithm updates the weights only in response to a mistake (i.e., $t_i \mathbf{x}_i^T \mathbf{w}^{k-1} \leq 0$ for some i). (Hint: Triangle inequality for the euclidean norm.)

Problem 9: Consider the cosine of the angle between $\tilde{\mathbf{w}}$ and \mathbf{w}^k and derive

$$k \leq \frac{R^2 \|\tilde{\mathbf{w}}\|^2}{\gamma^2}.$$

Now consider a new data set, \mathcal{D}' (again 2D inputs and two different classes):

$$\begin{aligned} \mathcal{D}' = & \{((0, 0), +1), ((-0.1, 0.1), +1), ((-0.3, -0.2), +1), ((0.2, 0.1), +1)\} \\ & \cup \{((0.2, -0.1), +1), ((-1.1, -1.0), -1), ((-1.3, -1.2), -1), ((-1, -1), -1)\} \\ & \cup \{((1, 1), -1), ((0.9, 1.2), -1), ((1.1, 1.0), -1)\} \end{aligned}$$

Problem 10: Can you separate this data with the perceptron algorithm? Why/why not?

Problem 11: Transform every input $\mathbf{x}_i \in \mathcal{D}'$ to \mathbf{x}'_i with $\mathbf{x}'_{i1} = \exp(\frac{-\|\mathbf{x}_i\|^2}{2})$ and $\mathbf{x}'_{i2} = \exp(\frac{-\|\mathbf{x}_i - (1, 1)\|^2}{2})$. If the labels stay the same, are the \mathbf{x}'_i s now linearly separable? Why/ why not?